


# BBM406

## Fundamentals of Machine Learning

### Lecture 5: ML Methodology





# About class projects



- This semester the theme is **machine learning for good**.
- To be done in groups of 3 people.
- **Deliverables:** Proposal, blog posts, progress report, project presentations (classroom + video presentations), final report and code
- For more details please check the project webpage:  
<http://web.cs.hacettepe.edu.tr/~aykut/classes/fall2019/bbm406/project.html>.

# Recall from last time... Linear Regression



$$y(x) = w_0 + w_1 x \quad \mathbf{w} = (w_0, w_1)$$

$$\ell(\mathbf{w}) = \sum_{n=1}^N \left[ t^{(n)} - (w_0 + w_1 x^{(n)}) \right]^2$$

**Gradient Descent Update Rule:**



$$\mathbf{w} \leftarrow \mathbf{w} + 2\lambda \left( t^{(n)} - y(x^{(n)}) \right) x^{(n)}$$

**Closed Form Solution:**

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$$

# Recall from last time... Some key concepts

- Data fits – is linear model best (**model selection**)?
  - Simplest models do not capture all the important variations (signal) in the data: **underfit**
  - More complex model may **overfit** the training data (fit not only the signal but also the **noise** in the data), especially if not enough data to constrain model
- One method of assessing fit:
  - test **generalization** = model's ability to predict the held out data




## Regularization

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

$$\|\mathbf{w}\|^2 \equiv \mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \dots + w_M^2$$

|         | $\ln \lambda = -\infty$ | $\ln \lambda = -18$ | $\ln \lambda = 0$ |
|---------|-------------------------|---------------------|-------------------|
| $w_0^*$ | 0.35                    | 0.35                | 0.13              |
| $w_1^*$ | 232.37                  | 4.74                | -0.05             |
| $w_2^*$ | -5321.83                | -0.77               | -0.06             |
| $w_3^*$ | 48568.31                | -31.97              | -0.05             |
| $w_4^*$ | -231639.30              | -3.89               | -0.03             |
| $w_5^*$ | 640042.26               | 55.28               | -0.02             |
| $w_6^*$ | -1061800.52             | 41.32               | -0.01             |
| $w_7^*$ | 1042400.18              | -45.95              | -0.00             |
| $w_8^*$ | -557682.99              | -91.53              | 0.00              |
| $w_9^*$ | 125201.43               | 72.68               | 0.01              |



# Today


- Machine Learning Methodology
  - validation
  - cross-validation (k-fold, leave-one-out)
  - model selection

# Machine Learning Methodology



# Recap: Regression

- In regression, labels  $y^i$  are continuous
- Classification/regression are solved very similarly
- Everything we have done so far transfers to classification with very minor changes
- Error: sum of distances from examples to the fitted model




# Training/Test Data Split

- Talked about splitting data in training/test sets
  - training data is used to fit parameters
  - test data is used to assess how classifier generalizes to new data
- What if classifier has “non-tunable” parameters?
  - a parameter is “non-tunable” if tuning (or training) it on the training data leads to overfitting
  - Examples:
    - k in kNN classifier
    - number of hidden units in MNN
    - number of hidden layers in MNN
    - etc ...




# Example of Overfitting

- Want to fit a polynomial machine  $f(\mathbf{x}, \mathbf{w})$
- Instead of fixing polynomial degree, make it parameter  $\mathbf{d}$ 
  - learning machine  $f(\mathbf{x}, \mathbf{w}, \mathbf{d})$
- Consider just three choices for  $\mathbf{d}$ 
  - degree 1
  - degree 2
  - degree 3
- Training error is a bad measure to choose  $\mathbf{d}$ 
  - degree 3 is the best according to the training error, but overfits the data



# Training/Test Data Split



- What about test error? Seems appropriate
  - **degree 2** is the best model according to the test error
- Except what do we report as the test error now?
- Test error should be computed on data that was **not used for training at all!**
- Here used “test” data for training, *i.e.* choosing model

# Validation data


- Same question when choosing among several classifiers
  - our polynomial degree example can be looked at as choosing among 3 classifiers (degree 1, 2, or 3)



# Validation data

- Same question when choosing among several classifiers
  - our polynomial degree example can be looked at as choosing among 3 classifiers (degree 1, 2, or 3)
- Solution: split the labeled data into three parts

labeled data



train tunable  
parameters  $w$

train other  
parameters,  
or to select  
classifier

use **only** to  
assess final  
performance

# Training/Validation

labeled data

Training  
 $\approx 60\%$

Validation  
 $\approx 20\%$


Test  
 $\approx 20\%$

Training error:  
computed on training  
example


Validation  
error:  
computed on  
validation  
examples

Test error:  
computed  
on  
test  
examples


# Training/Validation/Test Data



validation error: 3.3



validation error: 1.8




validation error: 3.4

- Training Data
- Validation Data
  - $d = 2$  is chosen
- Test Data
  - 1.3 test error computed for  $d = 2$




# Choosing Parameters: Example




- Need to choose number of hidden units for a MNN
  - The more hidden units, the better can fit training data
  - But at some point we overfit the data

# Diagnosing Underfitting/Overfitting




Underfitting

- large training error
- large validation error



Just Right

- small training error
- small validation error



Overfitting

- small training error
- large validation error

# Fixing Underfitting/Overfitting

- Fixing Underfitting
  - getting more training examples will not help
  - get more features
  - try more complex classifier
    - if using MLP, try more hidden units
- Fixing Overfitting
  - getting more training examples might help
  - try smaller set of features
  - Try less complex classifier
    - If using MLP, try less hidden units




# Train/Test/Validation Method

- Good news
  - Very simple
- Bad news:
  - Wastes data
    - in general, the more data we have, the better are the estimated parameters
    - we estimate parameters on 40% less data, since 20% removed for test and 20% for validation data
  - If we have a small dataset our test (validation) set might just be lucky or unlucky
- **Cross Validation is a method for performance evaluation that wastes less data**


# Small Dataset

Linear Model:




Mean Squared Error = 2.4

Quadratic Model:



Mean Squared Error = 0.9

Join the dots Model:




Mean Squared Error = 2.2

# LOOCV (Leave-one-out Cross Validation)

For  $k=1$  to  $n$

1. Let  $(\mathbf{x}^k, \mathbf{y}^k)$  be the  $k^{\text{th}}$  example







# LOOCV (Leave-one-out Cross Validation)

For  $k=1$  to  $n$

1. Let  $(\mathbf{x}^k, \mathbf{y}^k)$  be the  $k^{\text{th}}$  example
2. Temporarily remove  $(\mathbf{x}^k, \mathbf{y}^k)$  from the dataset




# LOOCV (Leave-one-out Cross Validation)



For  $k=1$  to  $n$

1. Let  $(\mathbf{x}^k, \mathbf{y}^k)$  be the  $k^{\text{th}}$  example
2. Temporarily remove  $(\mathbf{x}^k, \mathbf{y}^k)$  from the dataset
3. Train on the remaining  $n-1$  examples


# LOOCV (Leave-one-out Cross Validation)



For  $k=1$  to  $n$

1. Let  $(\mathbf{x}^k, \mathbf{y}^k)$  be the  $k^{\text{th}}$  example
2. Temporarily remove  $(\mathbf{x}^k, \mathbf{y}^k)$  from the dataset
3. Train on the remaining  $n-1$  examples
4. Note your error on  $(\mathbf{x}^k, \mathbf{y}^k)$

# LOOCV (Leave-one-out Cross Validation)




For  $k=1$  to  $n$

1. Let  $(\mathbf{x}^k, \mathbf{y}^k)$  be the  $k^{\text{th}}$  example
2. Temporarily remove  $(\mathbf{x}^k, \mathbf{y}^k)$  from the dataset
3. Train on the remaining  $n-1$  examples
4. Note your error on  $(\mathbf{x}^k, \mathbf{y}^k)$

When you've done all points, report the mean error

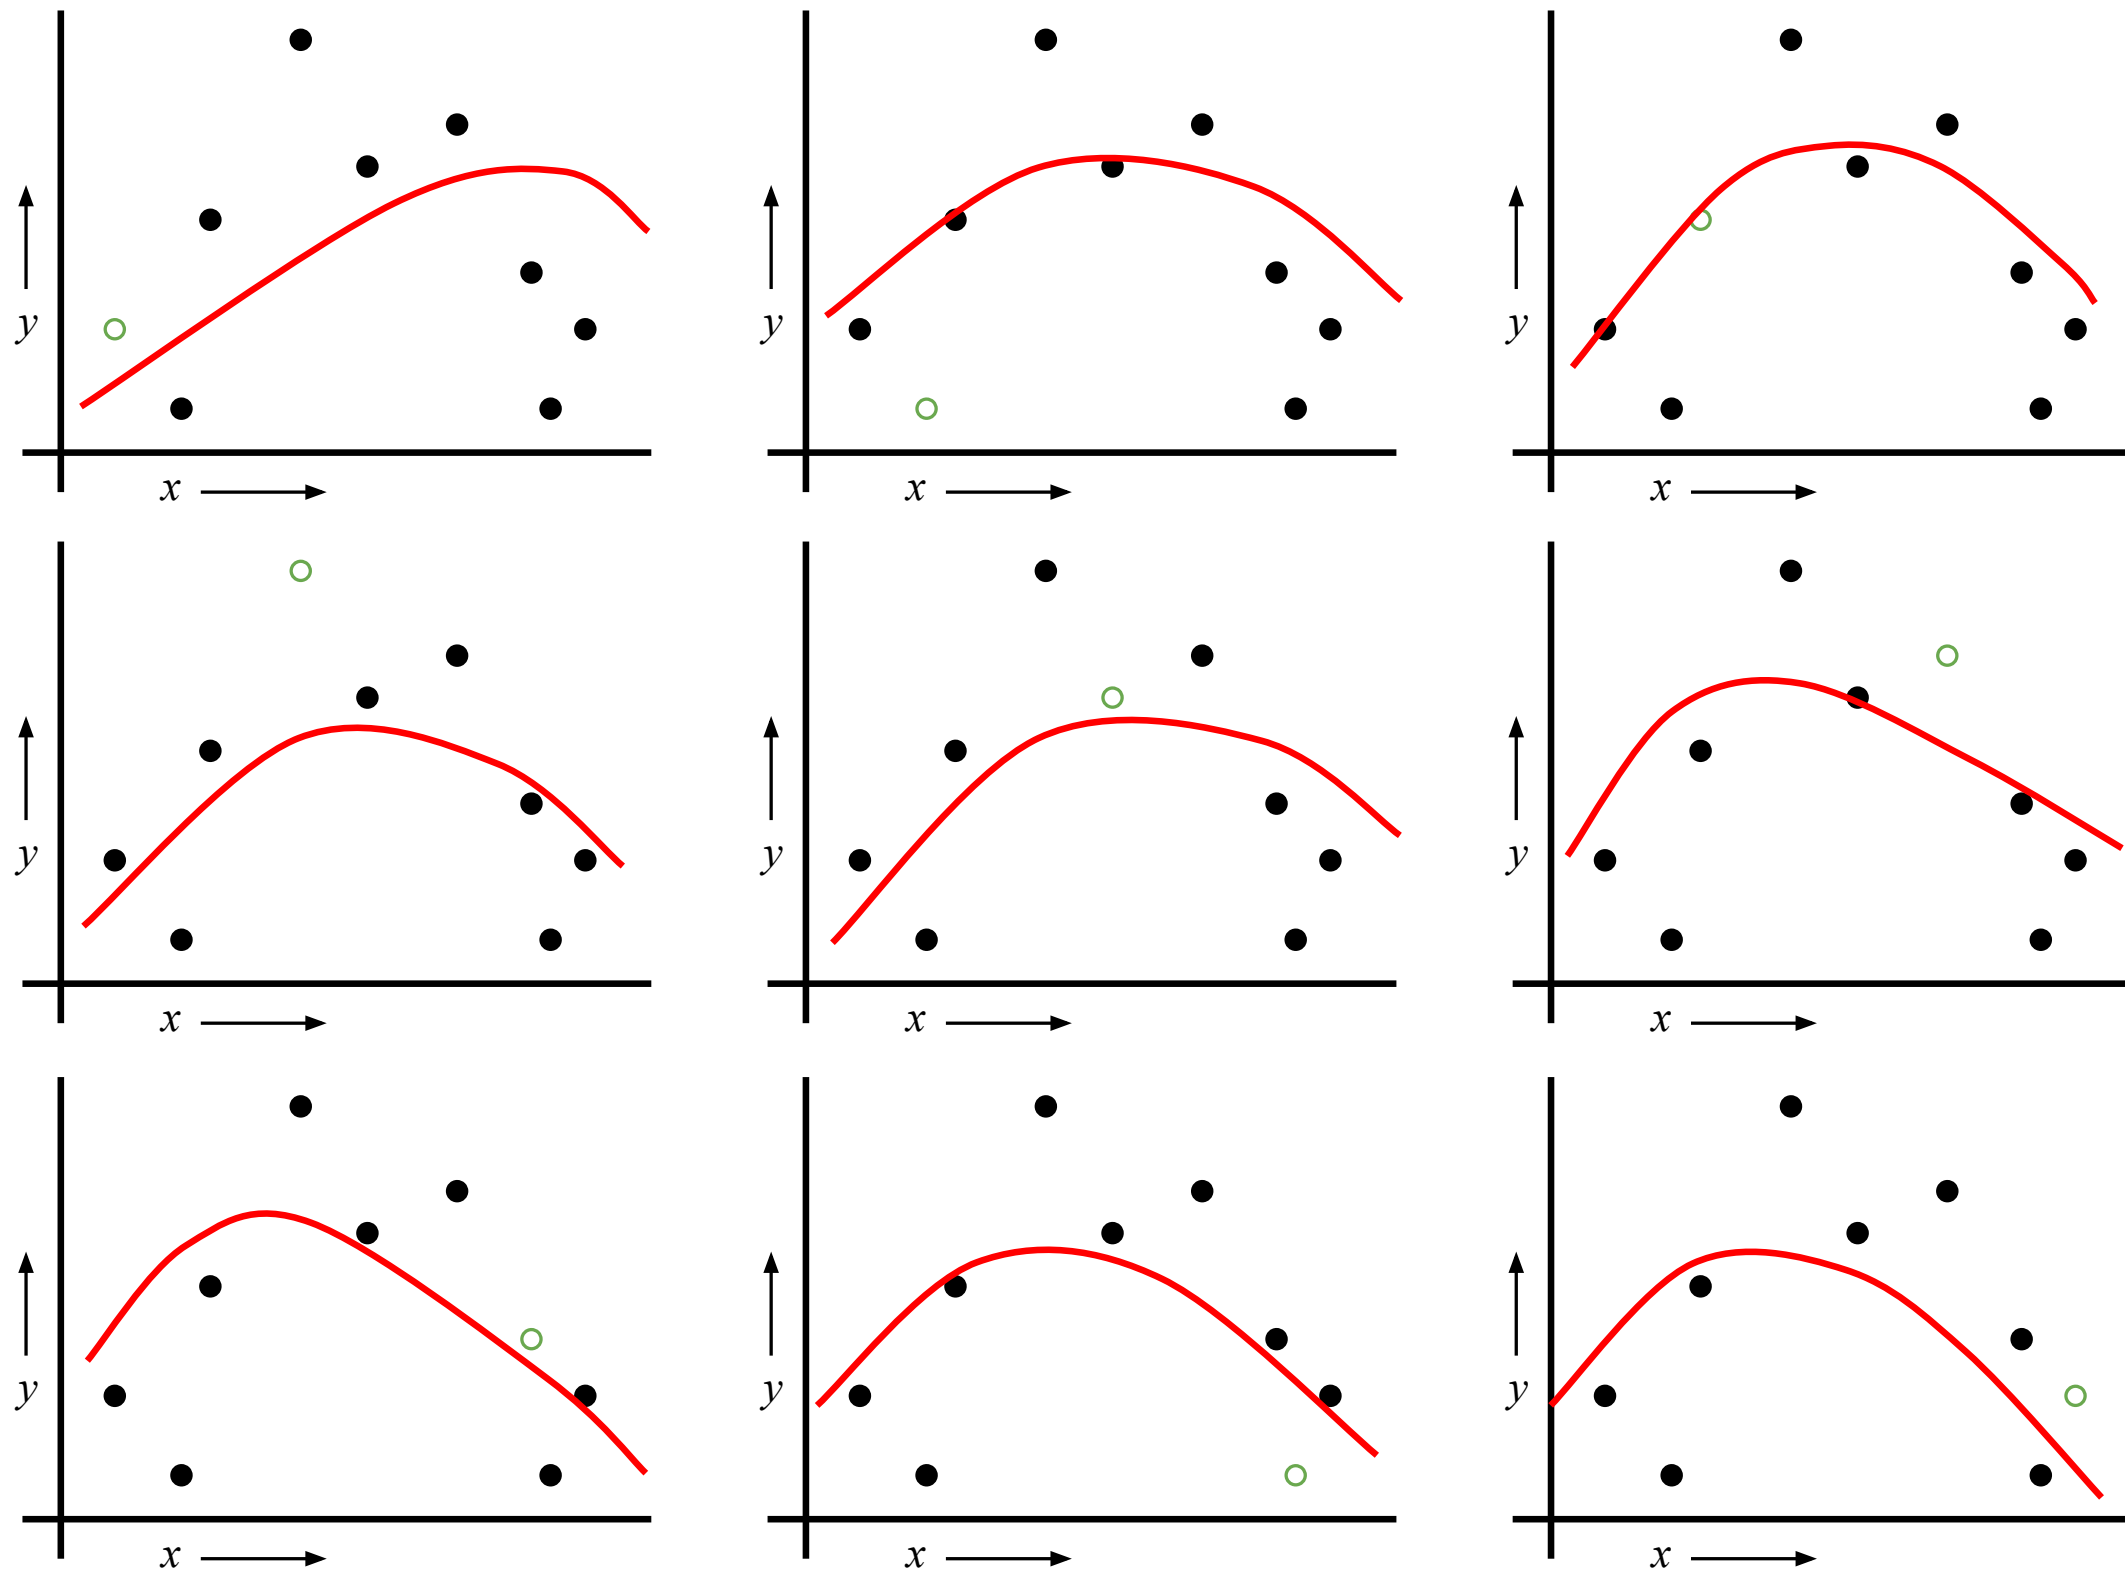


# LOOCV (Leave-one-out Cross Validation)



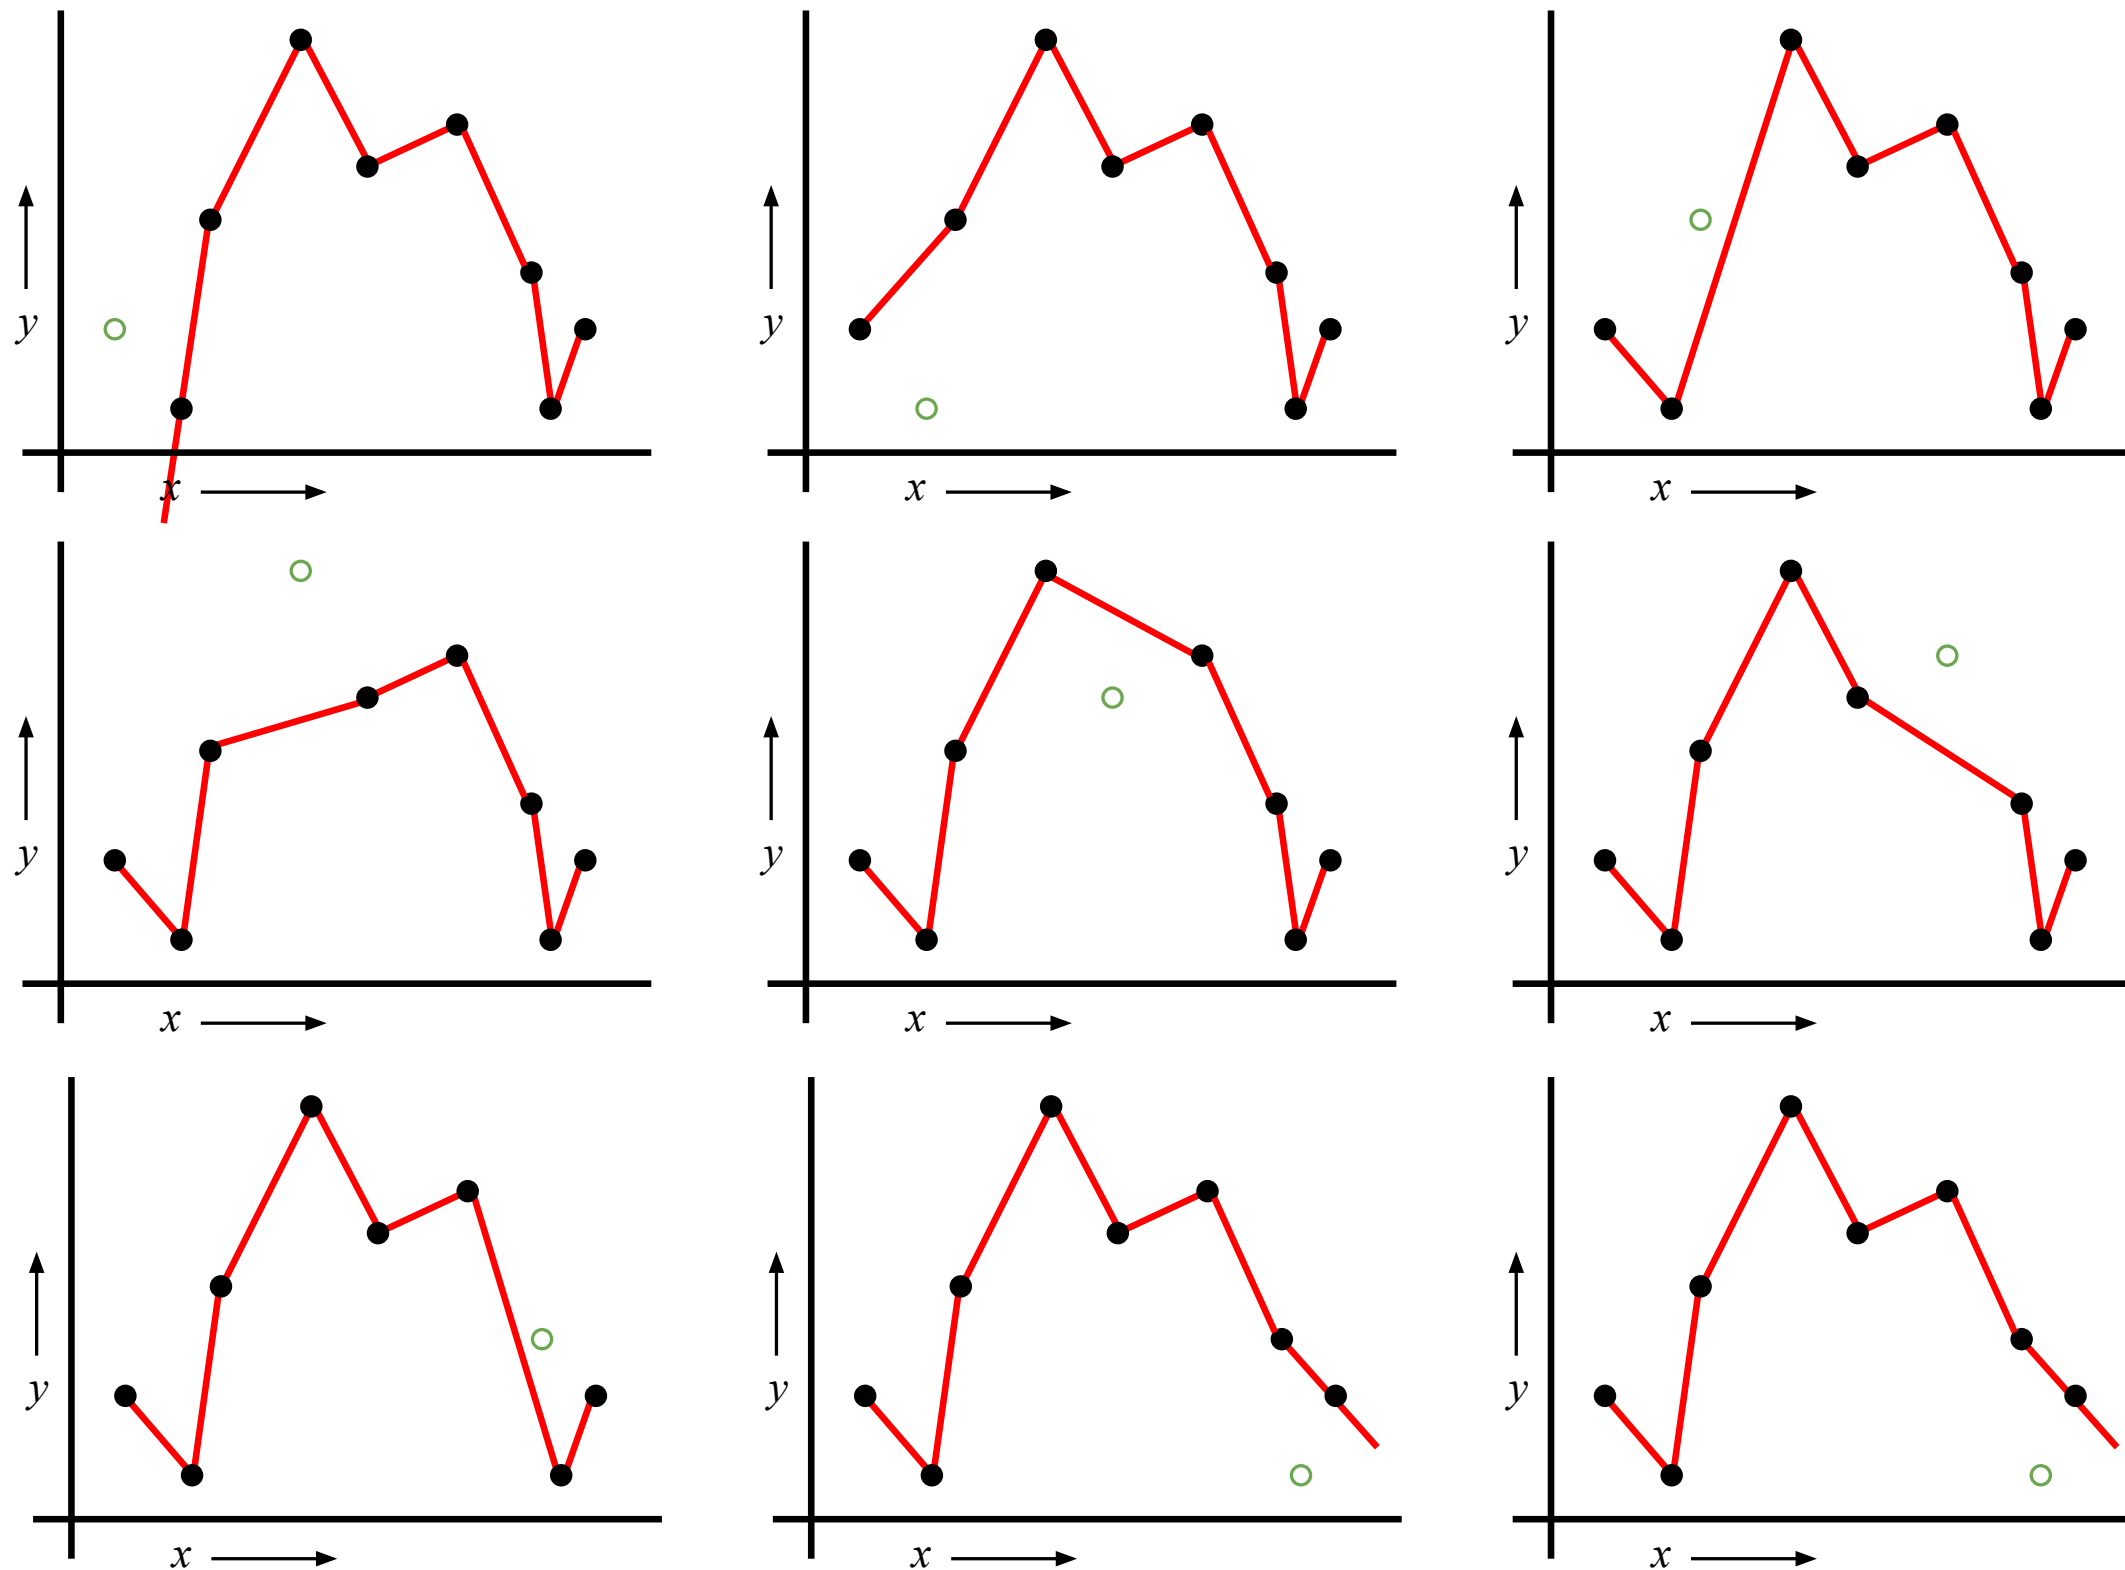
$$\text{MSE}_{\text{LOOCV}} = 2.12$$

# LOOCV for Quadratic Regression



$$\text{MSE}_{\text{LOOCV}} = 0.96$$

# LOOCV for Joint The Dots



$$\text{MSE}_{\text{LOOCV}} = 3.33$$


# Which kind of Cross Validation?

|                      | Downside   | Upside             |
|----------------------|--|--------------------|
| <b>Test-set</b>      | may give unreliable estimate of future performance | cheap              |
| <b>Leave-one-out</b> | expensive  | doesn't waste data |

- Can we get the best of both worlds?




# K-Fold Cross Validation




- Randomly break the dataset into  $k$  partitions
- In this example, we have  $k=3$  partitions colored red green and blue

# K-Fold Cross Validation




- Randomly break the dataset into  $k$  partitions
- In this example, we have  $k=3$  partitions colored red green and blue
- For the blue partition: train on all points not in the blue partition. Find test-set sum of errors on blue points

# K-Fold Cross Validation




- Randomly break the dataset into  $k$  partitions
- In this example, we have  $k=3$  partitions colored red green and blue
- For the blue partition: train on all points not in the blue partition. Find test-set sum of errors on blue points
- For the green partition: train on all points not in green partition. Find test-set sum of errors on green points

# K-Fold Cross Validation



- Randomly break the dataset into  $k$  partitions
- In this example, we have  $k=3$  partitions colored red green and blue
- For the blue partition: train on all points not in the blue partition. Find test-set sum of errors on blue points
- For the green partition: train on all points not in green partition. Find test-set sum of errors on green points
- For the red partition: train on all points not in red partition. Find the test-set sum of errors on red points

# K-Fold Cross Validation




Linear Regression

$MSE_{3FOLD} = 2.05$

- Randomly break the dataset into  $k$  partitions
- In this example, we have  $k=3$  partitions colored red green and blue
- For the blue partition: train on all points not in the blue partition. Find test-set sum of errors on blue points
- For the green partition: train on all points not in green partition. Find test-set sum of errors on green points
- For the red partition: train on all points not in red partition. Find the test-set sum of errors on red points
- Report the mean error



# K-Fold Cross Validation




Quadratic Regression

$$\text{MSE}_{3\text{FOLD}} = 1.1$$

- Randomly break the dataset into  $k$  partitions
- In this example, we have  $k=3$  partitions colored red green and blue
- For the blue partition: train on all points not in the blue partition. Find test-set sum of errors on blue points
- For the green partition: train on all points not in green partition. Find test-set sum of errors on green points
- For the red partition: train on all points not in red partition. Find the test-set sum of errors on red points
- Report the mean error

# K-Fold Cross Validation



Join the dots  
 $MSE_{3FOLD} = 2.93$

- Randomly break the dataset into k partitions
- In this example, we have k=3 partitions colored red green and blue
- For the blue partition: train on all points not in the blue partition. Find test-set sum of errors on blue points
- For the green partition: train on all points not in green partition. Find test-set sum of errors on green points
- For the red partition: train on all points not in red partition. Find the test-set sum of errors on red points
- Report the mean error

# Which kind of Cross Validation?

|               | Downside  | Upside   |
|---------------|---|--|
| Test-set      | may give unreliable estimate of future performance            | cheap  |
| Leave-one-out | expensive   | doesn't waste data   |
| 10-fold       | wastes 10% of the data, 10 times more expensive than test set | only wastes 10%, only 10 times more expensive instead of $n$ times |
| 3-fold        | wastes more data than 10-fold, more expensive than test set   | slightly better than test-set                                      |
| N-fold        | Identical to Leave-one-out                                    |  |

# Cross-validation for classification

- Instead of computing the sum squared errors on a test set, you should compute...

# Cross-validation for classification

- Instead of computing the sum squared errors on a test set, you should compute...


The total number of misclassifications on a test set

# Cross-validation for classification

- Instead of computing the sum squared errors on a test set, you should compute...

The total number of misclassifications on a test set

- What's LOOCV of 1-NN?
- What's LOOCV of 3-NN?
- What's LOOCV of 22-NN?











# Cross-validation for classification

- Choosing  $k$  for  $k$ -nearest neighbors
- Choosing Kernel parameters for SVM
- Any other “free” parameter of a classifier
- Choosing Features to use
- Choosing which classifier to use













# CV-based Model Selection

- We're trying to decide which algorithm to use.
- We train each machine and make a table...

| $f_i$ | Training Error  |
|-------|---|
| $f_1$ |   |
| $f_2$ |  |
| $f_3$ |  |
| $f_4$ |  |
| $f_5$ |  |
| $f_6$ |  |













# CV-based Model Selection

- We're trying to decide which algorithm to use.
- We train each machine and make a table...

| $f_i$ | Training Error  | 10-FOLD-CV Error  |
|-------|---|---|
| $f_1$ |   |   |
| $f_2$ |  |  |
| $f_3$ |  |  |
| $f_4$ |  |  |
| $f_5$ |  |  |
| $f_6$ |  |  |



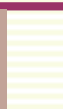









# CV-based Model Selection

- We're trying to decide which algorithm to use.
- We train each machine and make a table...

| $f_i$ | Training Error  | 10-FOLD-CV Error  | Choice |
|-------|---|---|--------|
| $f_1$ |   |   |        |
| $f_2$ |  |  |        |
| $f_3$ |  |  | ✓      |
| $f_4$ |  |  |        |
| $f_5$ |  |  |        |
| $f_6$ |  |  |        |

# CV-based Model Selection

- Example: Choosing “k” for a k-nearest-neighbor regression.
- Step 1: Compute LOOCV error for six different model classes:

| Algorithm | Training Error  | 10-fold-CV Error  | Choice |
|-----------|---|---|--------|
| k=1       |     |   |        |
| k=2       |    |  |        |
| k=3       |    |  |        |
| k=4       |    |  | ✓      |
| k=5       |   |  |        |
| k=6       |  |  |        |

- Step 2: Choose model that gave the best CV score
- Train with all the data, and that's the final model you'll use

# CV-based Model Selection

- Why stop at  $k=6$ ?
  - No good reason, except it looked like things were getting worse as  $K$  was increasing
- Are we guaranteed that a local optimum of  $K$  vs LOOCV will be the global optimum?
  - No, in fact the relationship can be very bumpy
- What should we do if we are depressed at the expense of doing LOOCV for  $k = 1$  through 1000?
  - Try:  $k=1, 2, 4, 8, 16, 32, 64, \dots, 1024$
  - Then do hillclimbing from an initial guess at  $k$

# **Next Lecture:**

## Learning Theory & Probability Review